

Deep learning for Images

...

ImageNet Project

ImageNet project: large visual database for recognition research

Ongoing research effort to provide easily accessible image database

Over 10 million images

Hand-annotated on what objects are pictured

ImageNet Project



ImageNet Categories: synsets





police van, police wagon,
paddy wagon, patrol
wagon, wagon, black
Maria



Komodo dragon, Komodo
lizard, dragon lizard, giant
lizard, Varanus
komodoensis

ImageNet Categories:



ILSVRC

The ImageNet Large Scale Visual Recognition Challenge (ILSVRC)

since 2010, annual software contest

Software programs compete to classify and detect objects and scenes

Model performance:

2011: error rate 25%

2012: CNN AlexNet 16%, runner up 26%

2015: error rate a few %, exceeds human ability at object recognition 1000 class

Convolutional Neural Nets (CNN)

Basic Units:

Input Layer

Convolutional layer


Relu layer

Pooling layer

FC layer

Input Layer

Input $[224 \times 224 \times 3]$: will hold the raw pixel values of the image, in this case an image of width 224, height 224, and with three color channels R,G,B.



Conv Layer

CONV layer will compute the output of neurons that are connected to local regions in the input, each computing a dot product between their weights and a small region they are connected to in the input volume.

This will result in an output volume of $[224 \times 224 \times 64]$ if we decided to use 64 filters, with each filter's receptive field size being 3×3 , using a stride of 1, and no zero-paddings.

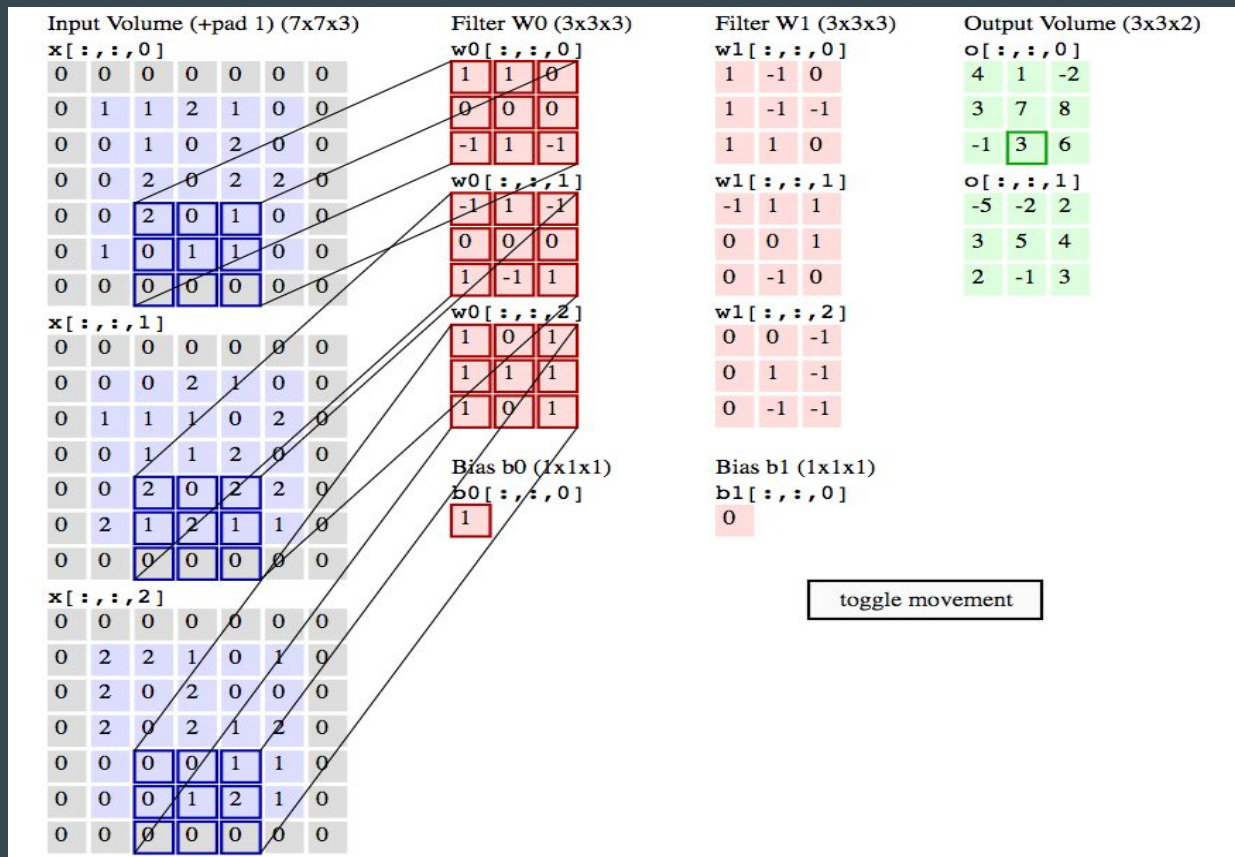
Conv Layer

CONV layer spatial arrangement: 3 hyperparameters control the size of output

- W: Input volume size
- F: Receptive Field size of the conv layer neurons
- S: Stride
- P: zero-padding on the border

Output size: $(W - F + 2P) / S + 1$

Conv Layer



$$W = 5, P = 1$$

$$F = 3, S = 2$$

Output size

$$= (5 - 3 + 2 \cdot 1) / 2 + 1$$

$$= 3$$

Output dimension:

$$3 \times 3 \times 2$$


Relu Layer

Rectified linear unit layer will apply an elementwise activation function, such as the $\max(0, x)$ thresholding at zero.

This leaves the size of the output volume unchanged ($[224 \times 224 \times 64]$).

Max Pooling

POOL layer will perform a downsampling operation along the spatial dimensions (width, height). If we use a max pooling of 2×2 filters with a stride of 2, the resulting output volume is $[112 \times 112 \times 64]$.



Fully Connected Layer

FC (i.e. fully-connected) layer: each neuron will be connected to all activations of the previous volume.

The output layer will compute class scores, resulting in volume size $[1 \times 1 \times 1000]$

ConvNets architecture summary

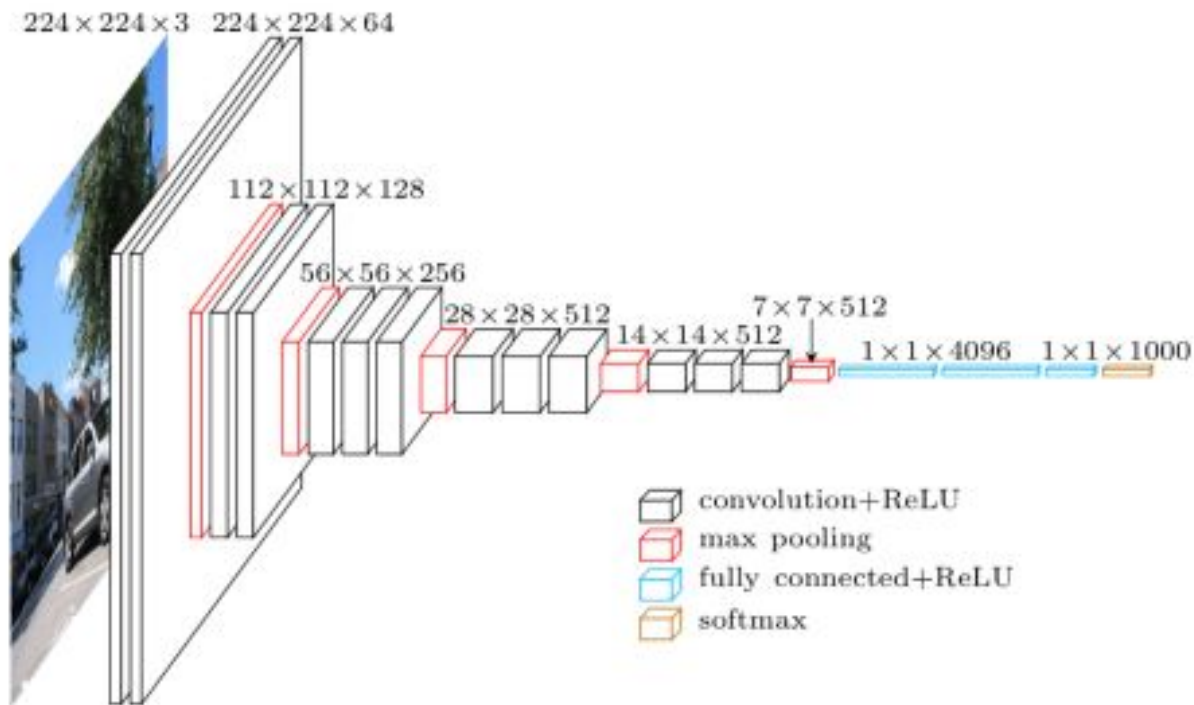
ConvNets transform the original image layer by layer from the original pixel values to the final class scores

Conv/FC layers perform transformations: weights and biases

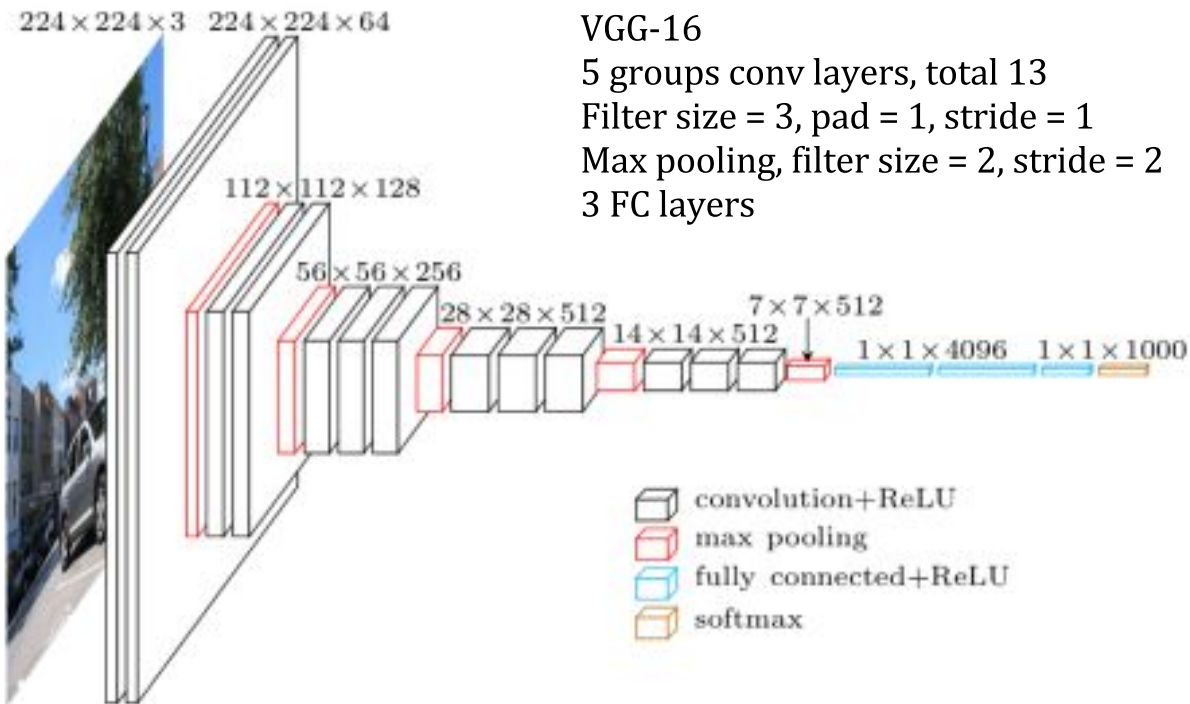
RELU/POOL layers will implement a fixed function

CONV/FC/POOL layers have additional hyperparameters (F, P, S), RELU don't


ConvNet Example in Keras!



ConvNet Example in Keras !



ConvNet Example in Keras !




ConvNet Example in Keras !


```
model = Sequential()  
model.add(ZeroPadding2D((1,1),input_shape=(3,224,224)))  
model.add(Convolution2D(64, 3, 3, activation='relu'))  
model.add(ZeroPadding2D((1,1)))  
model.add(Convolution2D(64, 3, 3, activation='relu'))  
model.add(MaxPooling2D((2,2), strides=(2,2)))  
  
model.add(ZeroPadding2D((1,1)))  
model.add(Convolution2D(128, 3, 3, activation='relu'))  
model.add(ZeroPadding2D((1,1)))  
model.add(Convolution2D(128, 3, 3, activation='relu'))  
model.add(MaxPooling2D((2,2), strides=(2,2)))
```

ConvNet Example in Keras !


conv1_1: a few of the 64 filters



conv2_1: a few of the 128 filters



ConvNet Example in Keras !



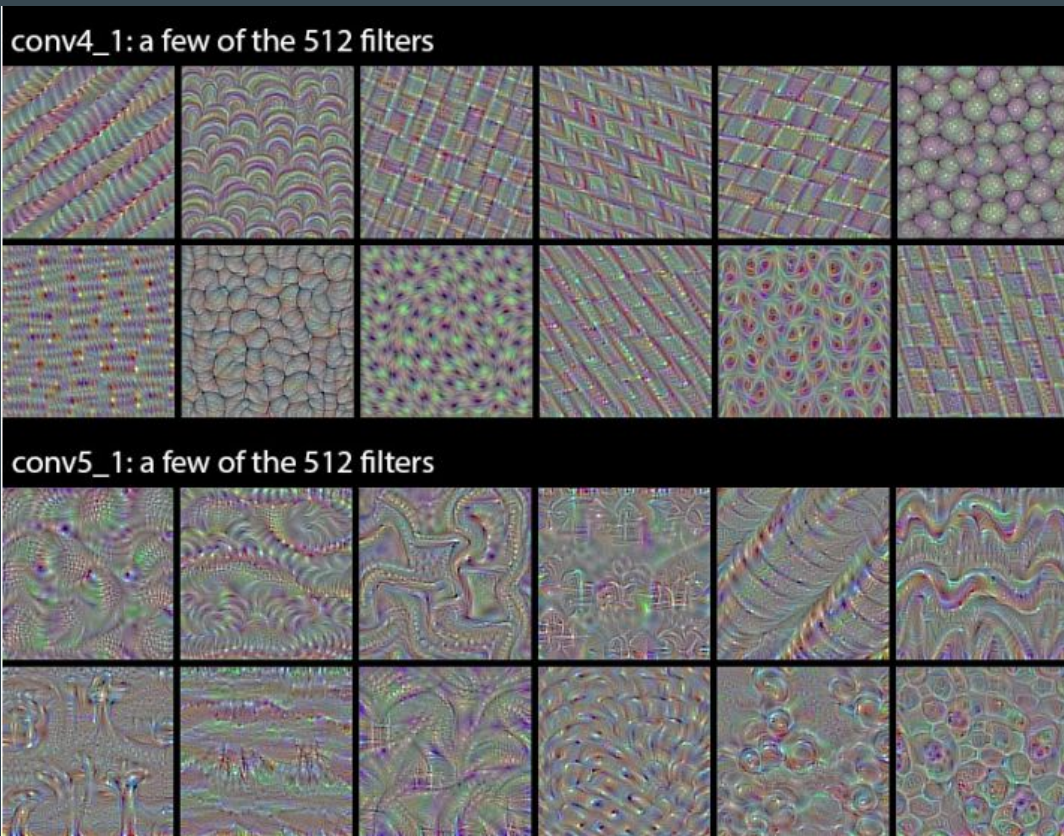
ConvNet Example in Keras !

```
model.add(ZeroPadding2D((1,1)))
model.add(Convolution2D(256, 3, 3, activation='relu'))
model.add(ZeroPadding2D((1,1)))
model.add(Convolution2D(256, 3, 3, activation='relu'))
model.add(ZeroPadding2D((1,1)))
model.add(Convolution2D(256, 3, 3, activation='relu'))
model.add(MaxPooling2D((2,2), strides=(2,2)))
```


```
model.add(ZeroPadding2D((1,1)))
model.add(Convolution2D(512, 3, 3, activation='relu'))
model.add(ZeroPadding2D((1,1)))
model.add(Convolution2D(512, 3, 3, activation='relu'))
model.add(ZeroPadding2D((1,1)))
model.add(Convolution2D(512, 3, 3, activation='relu'))
model.add(MaxPooling2D((2,2), strides=(2,2)))
```

```
model.add(ZeroPadding2D((1,1)))
model.add(Convolution2D(512, 3, 3, activation='relu'))
model.add(ZeroPadding2D((1,1)))
model.add(Convolution2D(512, 3, 3, activation='relu'))
model.add(ZeroPadding2D((1,1)))
model.add(Convolution2D(512, 3, 3, activation='relu'))
model.add(MaxPooling2D((2,2), strides=(2,2)))
```

ConvNet Example in Keras !



ConvNet Example in Keras!



ConvNet Example in Keras !

```
model.add(Flatten())  
model.add(Dense(4096, activation='relu'))  
model.add(Dropout(0.5))  
model.add(Dense(4096, activation='relu'))  
model.add(Dropout(0.5))  
model.add(Dense(1000, activation='softmax'))
```


ConvNet Example in Keras !

